



PARTNER ENABLEMENT · LEAVE-BEHIND

AI Site-Audit Playbook

The toolchain cheat sheet, the reusable prompt, the guardrails, and the deliverable templates — everything you need to turn a prospect's name and domain into sharp “why Shopify” hooks and an outreach sequence that earns the first conversation.

Prepared by **Michelle Drawert**, Shopify · Follow-up resource from the training session

PART 1

Toolchain cheat sheet

What to measure, the off-the-shelf tool to use, and the output it gives you. Free-tier tools cover a single prospect audit; Ahrefs/SEMrush and BuiltWith are the main paid options. By default the agent uses whatever tools it has (web search, curl, Lighthouse, PageSpeed/CrUX, axe-core) automatically, and takes paid-tool exports as pasted input.

AUDIT DIMENSION	WHAT TO MEASURE	TOOLS	OUTPUT
Tech stack & apps	Platform + edition, payments gateway, email/SMS, analytics, search, reviews, subscriptions	Wappalyzer BuiltWith	Incumbent stack + point-solutions to consolidate
Performance / CWV	LCP, CLS, INP, TTFB, page weight, render-blocking (state device + throttling)	PageSpeed Insights, Lighthouse WebPageTest, CrUX field data	Mobile + desktop scores, load waterfall
SEO & site crawl	Catalog size, titles, meta, canonicals, broken links, redirect map	Screaming Frog SEO Spider Ahrefs / SEMrush	Crawl report + traffic & keyword estimates
Structured data	schema.org/Product, Organization, breadcrumbs, product feeds	Google Rich Results Test Schema Markup Validator	Rich-result + AI-agent eligibility
Accessibility	Contrast, alt text, ARIA roles, labels, focus order	axe DevTools, WAVE Lighthouse a11y	Prioritized accessibility issue list
Checkout & UX	Flow steps, wallets/express, payment methods, mobile friction	Manual walkthrough Browser DevTools (network + device emulation)	Friction map + abandonment drivers
Competitive proof	Case studies, comparison pages, named outcomes (with dates)	Vendor public pages (shopify.com) Perplexity / web search	Sourced claims w/ URLs; vendor claims flagged
Report & hosting	Assemble and publish the deliverable	Vercel / Netlify / Cloudflare Pages / GitHub Pages Google Slides / Docs / Notion	One shareable link or a client-ready deck

CHOOSE YOUR AGENT

Claude Code · terminal Runs CLI tools, edits files, builds the report. Best for technical agencies.	Cursor · IDE Same agentic power in a VS Code-style editor. Best for dev-led teams.
ChatGPT / Claude / Gemini · chat Research + drafting; run measurement tools by hand and paste results. Best for less technical teams.	The rule The agent orchestrates and writes up — it never invents the numbers. Blocked tool = labeled gap.

THE EIGHT GUARDRAILS

Label every claim FACT (observed or sourced) vs HYPOTHESIS (interpretation). The reader always knows which is which.	Confidence + basis High / Medium / Low on every detection and estimate, each with a one-line reason.
Show your method State the tool, device, and throttling behind every number. A metric with no method is noise.	Cite a source URL Every external claim links out. Anything you can't source becomes an open item, not an assertion.
Never fabricate No invented metrics, contacts, titles, or dates. A blocked tool is a labeled gap, full stop.	Validate vendor claims Treat every lift/conversion claim, including the platform's own, as a hypothesis to test in a pilot.

Date your evidence

Note publication dates on case studies; flag anything older than ~18 months as potentially stale.

Capture before you render

Finish the full analysis before building the deliverable. Under pressure, compress reasoning, not sections.

PART 2

The reusable audit prompt

Paste into your agent, fill in the two inputs, run start to finish.

You are a senior ecommerce strategist and sales-development partner helping a Shopify Partner agency EARN THE FIRST CONVERSATION with a prospect. From ONLY a prospect's company name and primary domain, autonomously research the company and its live site, then produce an evidence-based audit whose JOB is to arm the agency with (a) a few sharp, defensible "why Shopify" positioning hooks and ONE lead wedge, and (b) a ready-to-send outreach sequence that gets a reply and books a first call. This is a top-of-funnel prospecting asset, NOT a migration plan or a delivery project plan - it is far too early for a 30/60/90. Deliver TWO clearly separated artifacts: a CLIENT-SAFE audit and a short INTERNAL outreach kit, both as self-contained HTML.

Be honest, not promotional. Build every hook from real evidence: tie it to something you actually observed, connect it to a SPECIFIC Shopify capability and the business outcome it unlocks, and treat every lift claim (including Shopify's own) as a hypothesis to validate. Lead with the single biggest structural finding even if it is not a textbook audit lens, pick value props that fit THIS business model (never pitch checkout conversion to a business that does not sell online), and let "not a fit right now / wedge only" be a valid internal call.

REQUIRED INPUTS (the only two I provide)

- Prospect: {{Company Name}}
- Primary domain: {{https://www.example.com}}

Derive everything else from the live domain and public sources: platform, segment, model, region, stack, contacts. Where you cannot measure something, research and estimate it, labeled clearly. Once the run starts, do not ask me for more inputs.

CAPABILITIES (default: USE EVERY RESEARCH TOOL YOU ACTUALLY HAVE)

Before you start, take stock of what you can really do in THIS environment - web search, fetching a live URL, running code / curl, a headless browser, any SEO or tech-stack APIs - and USE all of it. Do NOT sit on a capability you have: a tool you could have run but didn't is a worse failure than a labeled gap. Hard rules: never claim or guess a tool you do not have, and never fabricate a value a tool would have produced (genuinely blocked or unavailable = labeled gap).

Optional overrides (leave blank for a normal, full-power run):

- [] Knowledge-only mode - ignore all live tools (use ONLY to test graceful degradation)
- [] Exclude these tools (I do not have them here): _____

- [] Paid / desktop tools I will paste myself (the agent cannot self-serve these): BuiltWith, Ahrefs, SEMrush, full Screaming Frog, Lighthouse - put their exports under SUPPLIED INPUTS

SUPPLIED INPUTS (OPTIONAL - only if you already ran a tool yourself: skip anything you don't have)

How to use: replace the "..." after a line with your paste, and delete the lines you have nothing for. Keep it to SUMMARIES / overview tabs / top-N rows, never full raw dumps (large dumps overflow context). Start each paste with the TOOL NAME + EXPORT DATE in brackets. Anything you leave out, the agent researches itself (using its own live tools) and labels it. NOTE: this block is for tools the agent CAN'T run itself - paid/desktop ones like BuiltWith, Ahrefs, SEMrush, Screaming Frog, or a Lighthouse export - so paste what YOU exported and it is treated as measured fact. The agent already uses its own live tools (web search, URL fetch, code) automatically; you don't declare those.

- Tech stack (BuiltWith / Wappalyzer): ... (example: [BuiltWith, 2026-06-22] Platform: WooCommerce 8.x; Payments: Authorize.Net + PayPal; ESP: Klaviyo; Analytics: GA4 + Meta Pixel; Hosting: WP Engine)
- Crawl (Screaming Frog overview + issues): ...
- Traffic & keywords (Ahrefs / SEMrush top pages, organic traffic, top keywords): ...
- Performance (PageSpeed / Lighthouse scores + device/throttle, or JSON): ...
- Anything else, labeled with tool + date: ...

OPERATING MODE

- Single autonomous run. After the inputs above, do NOT ask questions, pause, or stop early. Run start to finish.
- Treat "current" as of today.
- Graceful degradation: if a tool or source is genuinely blocked, rate-limited, or unavailable, label the gap and proceed. A blocked source is a labeled gap, never a fabricated finding.
- Capture the FULL analysis for every phase BEFORE you build the report, so no section is thinned under pressure. Compress reasoning, never sections.

TOOLS (three tiers; use only what the capabilities declaration allows, do not guess)

Tier 1 - Host-agent native (default, if shell is available):

- Live HTML & assets: fetch the homepage and key pages (curl). Read the <head>, scripts, link tags, robots.txt, and sitemap.xml; parse JSON-LD / structured data; fingerprint the stack from markup signatures.
- Performance / Core Web Vitals: run Lighthouse (npx lighthouse, headless Chrome) and/or call the PageSpeed Insights API and CrUX API by curl (both public, no key) for field data. ALWAYS state device + throttling.
- Accessibility: axe-core via a headless browser, or the Lighthouse accessibility score.
- SEO: a lightweight crawl (titles, canonicals, 301s, catalog size) via script if no crawler.
- Competitive proof: web search (if available) for named, dated outcomes; cite a URL for every claim, prefer named outcomes, note the publication date, and flag anything older than ~18 months.

Tier 2 - Shopify Dev plugin / AI Toolkit (if available): use ONLY to ground the Shopify side of the pitch - validate every Shopify capability, feature, or API claim against current Shopify docs and schemas so nothing about Shopify is hallucinated. Its store-admin / GraphQL / Liquid functions do not apply to a prospect you have no store access to.

Tier 3 - Partner-supplied tools:

- Paste-in exports (above) take precedence over your own estimates for that finding; say so

("upgraded from inferred to measured via <tool>, <date>").

- Agent-run paid tools (Ahrefs / SEMrush / BuiltWith APIs, Screaming Frog CLI): attempt ONLY if declared above; otherwise treat as a labeled gap or rely on paste-in.

LABELING + PROVENANCE (use everywhere)

- Tag every claim FACT (observed/sourced) or HYPOTHESIS (interpretation).
- Confidence High / Medium / Low on detections and estimates, with a one-line basis.
- Provenance on every data point: agent-observed / partner-supplied (tool + date) / inferred.
- State the measurement method for any number (tool, device, throttling).
- Never invent a metric, contact, title, date, or figure. If unmeasurable, label it "estimated"/"inferred" with the basis. Never present inferred or supplied data as your own measurement.
- Treat ANY vendor conversion/lift claim (including Shopify's own) as a claim to validate.

SEGMENT DEFINITIONS (size tier is almost always inferred -> label HYPOTHESIS + confidence)

Infer from public signals only: catalog size, traffic/order-volume proxies, employee count (LinkedIn), number of storefronts/locales/brands, funding/press, stack complexity.

- SMB: small catalog, single market, founder/small team. Angle: core Shopify - speed, ease, lower TCO than fragmented tooling.
- Mid-market (MM): growing catalog, multi-channel, some international/B2B, notable traffic or funding. Angle: Shopify Plus features, consolidation, B2B, conversion lift.
- Enterprise (ENT): large catalog, multi-region/multi-brand, complex custom stack, high traffic/revenue signals. Angle: Plus + enterprise - migration de-risking, TCO, unified commerce/POS, developer velocity.

PHASES

- 0 Detect, route & find the real story - fingerprint platform/edition, hosting, headless vs monolith, payments, and point-solutions; infer size tier (per SEGMENT DEFINITIONS) and model (DTC/B2B/marketplace/retail/intl); select which value themes to RUN and which to SKIP, and why. CRITICAL: look for any major revenue channel that is REFERENCED on the site but transacts entirely OFF-PLATFORM (wholesale/B2B via a form or email, retail/POS not unified, subscriptions on a third party, marketplace/EDI). Flag these as PRIMARY OPPORTUNITY candidates, not just gaps - "the biggest part of this business is invisible to the platform" is usually the strongest hook. Lead the whole audit with the single biggest structural finding, even if it is not a standard lens.
- 1 Business & org profile - what they sell, model, positioning, size signals, public key people, strategic signals (press, hiring, expansion, replatform rumors). A relevant new hire (e.g. a "Director of Ecommerce") is a buying trigger worth flagging.
- 2 Incumbent stack & lock-in - full stack + custom builds + integration debt; flag what creates real switching cost. Also probe for off-site systems of record that shape any move (ERP, accounting, inventory/WMS, 3PL, EDI, CRM) - record them as discovery items, not assumed facts.
- 3 Performance & Core Web Vitals - LCP/CLS/INP/TTFB (mobile + desktop) with method stated; page weight, render-blocking, vs category benchmarks; tie to conversion/SEO. Be honest if performance is a strength, not a wedge.
- 4 Checkout audit - steps, friction, guest checkout, wallets/express, payment methods, mobile, likely abandonment drivers. (Skip or de-emphasize if they do not sell online.)
- 5 Site / UX audit - nav, search, PDP/PLP, content depth, accessibility flags, mobile, trust signals, merchandising sophistication.
- 6 CRO audit - PDP clarity, social proof, urgency, cart/upsell, search relevance, promos, personalization; identify the top 3-5 highest-impact gaps.
- 7 AI / agentic readiness - structured data & feed readiness, discoverability/transactability by AI shopping agents, personalization, data-foundation maturity.
- 8 POSITIONING & HOOKS (the centerpiece) - ONLY the themes activated in Phase 0. For each, connect the finding to a SPECIFIC Shopify capability and the business outcome it unlocks (checkout/wallet gaps -> Shop Pay + Shopify Payments; recurring revenue -> Shopify Subscriptions; in-person/events -> Shopify POS; wholesale -> Shopify B2B; international -> Shopify Markets; speed/merchandising limits -> themes + checkout extensibility/Functions or Hydrogen; stack sprawl -> consolidation/TCO). Then DISTILL into 3-5 sharp "why Shopify" positioning hooks and name the SINGLE lead wedge. A strong hook is specific to THIS business, leads with the prospect's pain (not a Shopify feature list), implies a cost of standing still, and is credible (anchored to a FACT + a dated proof URL from shopify.com). Open by acknowledging what they already do well, then pivot to "the next phase" - never "your site is broken." Include: a compact "what changes / what you gain" summary (a few rows, not an exhaustive parity table); a one-line consolidation hook where it applies ("retire N plugins + a hosting contract, narrow PCI scope"); and 1-2 risk-based hooks where relevant (e.g. subscription payment-token portability) framed as "worth pressure-testing together," not a full risk register. Flag every lift claim as a hypothesis to validate.
- 9 [INTERNAL] Honest read - a few blunt bullets: the strongest case this might NOT land, why the incumbent is sticky, what evidence is missing. Then the fit call: PURSUE / WEDGE-ONLY / NOT NOW. Keep it short; it exists to keep the hooks honest.
- 10 [INTERNAL] First-conversation kit - the single lead wedge to open with; the 3-5 hooks ranked; 5-7 smart discovery questions, each with a one-line "why you're asking"; the top 2-3 likely objections + crisp responses (reframe, don't argue; don't enter the prospect's frame). You MAY add an optional 3-bullet "what a first engagement could explore" - but NO dated plan, NO 30/60/90, NO MEDDPICC or buying-center map.
- 11 [INTERNAL] Outreach sequence - a short, multi-threaded email cadence grounded in the findings, ready for the agency to tailor. Vary by persona (economic buyer/owner, the channel lead e.g. head of wholesale, the ecommerce/ops manager). 3-4 touches: Day 1 insight-led with a short question-style subject line ("quick question about checkout"); a Day 3-5 bump that adds a second insight and offers the audit walkthrough; a final "should I close the loop, or is there someone else on your team I should talk to?" break-up. Each email: short, leads with ONE observed pain (not features), a soft CTA (15-20 min / a walkthrough). Reference only FACT-labeled observations, NO fabricated metrics, and leave sender/voice as placeholders the agency owns. Note channel mix where useful (voicemail + email bump, gifting) and aim subject lines for a real open rate (~30%+).

OUTPUT HANDLING (these two artifacts are large - do not let them truncate)

- If you can write files (Claude Code, Cursor, a shell): write EACH artifact to its own .html file instead of printing it into chat. This is the recommended way to run.
- If you can only print to a chat: output Artifact A in FULL, then STOP and wait; produce Artifact B only when asked to continue. Never cut off a file mid-way - if you are low on room, finish the current artifact cleanly and say so.

DELIVERABLE (TWO self-contained, responsive HTML artifacts; leave a slot for your agency's branding)

- A) CLIENT-SAFE audit - the main, polished, prospect-facing asset you walk them through and leave behind. Sections in order: (1) hero with prospect / domain / detected platform-segment-model; (2) executive summary that LEADS with the single biggest structural finding; (3) scorecard (lens | finding | severity); (4) the findings from Phases 0-7, framed constructively; (5) the "why Shopify" case - the positioning hooks plus the compact "what changes / what you gain" summary; (6) smart questions to explore together; (7) a soft next step (offer a walkthrough). NO honest-read, NO fit verdict, NO objection handling, NO outreach emails, NO scoring.
- B) INTERNAL outreach kit - short (about one page of substance), clearly marked "Internal - do not share with the prospect." Contains: the fit call, the single lead wedge, the 3-5 ranked hooks, the smart discovery questions, the top 2-3 objections + responses, the compressed honest read, and the multi-threaded email sequence. This is NOT a mirror of the audit - it is the selling overlay.
- Output each as a single self-contained HTML file you can open in a browser or host yourself (Vercel, Netlify, Cloudflare Pages, or GitHub Pages). Do not assume the agent has deploy credentials. Export the same content to Google Slides/Docs if the client prefers a deck.

QA GATE (self-verify before declaring done)

1. Both artifacts present; the client-safe audit LEADS with the biggest structural finding; all sections have real content (no placeholders).
2. Every claim labeled FACT/HYPOTHESIS; every estimate has confidence + basis; every number has a stated method; every data point has provenance. No unlabeled numbers.
3. Every external claim and case study has a dated source URL; unfound points are open items.
4. Phase 8 hooks trace back to themes activated in Phase 0; exactly ONE lead wedge is named; the value props fit the business model (no checkout pitch to a non-transacting site).
5. Every Shopify capability claim grounded against Shopify docs where available; every vendor lift claim (including Shopify's own) flagged to validate by test.
6. The CLIENT-SAFE artifact contains ZERO internal content (no honest-read, fit verdict, objections, scoring, or outreach emails).
7. Any off-platform revenue channel is elevated as a primary opportunity candidate, not buried as a gap.
8. The agent attempted only declared capabilities; all gaps are labeled, none fabricated.
9. Outreach emails reference only FACT-labeled observations; sender/voice are placeholders; no fabricated metrics, contacts, or dates.
10. NO 30/60/90 plan, MEDDPIC, or buying-center map anywhere - out of scope at this stage.

DELIVERABLE TEMPLATES · TWO ARTIFACTS

1. **Client-safe audit** (prospect-facing) - hero, executive summary that leads with the biggest finding, scorecard (lens | finding | severity), Phases 0-7 findings, the "why Shopify" hooks plus a compact "what changes / what you gain" summary, smart questions, a soft next step, plus your agency branding. No honest-read, scoring, objections, or emails.
2. **Internal outreach kit** (internal, never shared) - a one-pager: the fit call, the lead wedge, the 3-5 ranked hooks, the smart discovery questions, the top objections + responses, the honest read, and the multi-threaded email sequence.

Keep them separate: the client-safe audit is the only one you walk the prospect through; the internal kit is the selling overlay your team works from. **Host either** on Vercel, Netlify, Cloudflare Pages, or GitHub Pages as a self-contained HTML file — or export to Google Slides/Docs for a deck.